

Class Imbalance in Histopathology Patch Classification

Yannik Queisler

September 25, 2026

1 Motivation

Histopathology datasets are long-tailed almost by construction. Common cancers and benign tissue dominate archives, while rare subtypes and atypical lesions, often the diagnostically decisive ones, contribute few patients and patches. BRACS, for example, was released with explicitly unbalanced atypical classes (Brancati et al., 2022). In CNNs, imbalance biases predictions toward frequent classes and lowers minority recall (Buda et al., 2018); in medical imaging it also distorts calibration, and conclusions depend on whether evaluation uses balanced metrics (Mosquera et al., 2024). Long-tailed medical benchmarks show that no single remedy works across datasets (Ju et al., 2024), and histopathology-specific models keep being proposed (Cong et al., 2022). Frozen foundation-model features with a linear probe are now the standard baseline (Chen et al., 2024; Zimmermann et al., 2024), yet how much imbalance still costs such a probe, and why, has not been sufficiently measured in isolation.

2 Setup

A multinomial logistic regression is trained on frozen Virchow2 patch features (Zimmermann et al., 2024) on two tasks: TCGA-UT, 30 cancer types (Komura et al., 2022), and BRACS, 7 breast-lesion subtypes from normal tissue to invasive carcinoma (Brancati et al., 2022). Every condition uses the same patients per class (20 on TCGA-UT, 10 on BRACS) and the same total patch budget; only the share of that budget each class receives changes. An exponential profile sets the ratio ρ between the largest and smallest class to 1–100 (Figure 1), class-to-rank assignment is randomized over 30 draws (3 splits $\times$ 10), and a native condition keeps each dataset’s own class shares. The outcome is the *damage* $\Delta(\rho) = \text{BA}(1) - \text{BA}(\rho)$, the loss in test patient-macro balanced accuracy against the balanced condition, with paired bootstrap 95% intervals; damage counts as material from 1 percentage point (pp).




Figure 1: Training class distributions, classes sorted from head to tail: the balanced condition, the native class shares of one split, and the most imbalanced ratio condition. The tested ratios are $\rho \in \{1, 2, 5, 10, 20, 50, 100\}$; the profiles for $\rho = 2$ to 50 lie between the balanced and the $\rho = 100$ profile. Which class sits at which rank is randomized per draw.




Figure 2: (a) Balanced-accuracy damage versus ρ with 95% bands; diamonds mark native class shares, dotted line the 1-pp threshold. (b) Expected calibration error (ECE) of the raw predicted probabilities (solid) and after a single global confidence correction fitted on validation data (dashed).

3 Effects of class imbalance

Balanced accuracy reaches 72.1% on TCGA-UT and 39.6% on BRACS. Published evaluations show the same broad gap: ROI-level linear probes score 60–68% on TCGA-UT (Mahmood Lab, 2025), while slide-level linear probes score 34.8–48.1% on BRACS (Vaidya et al., 2025). TCGA-UT spans cancer types across organs, while BRACS separates breast lesions with similar morphology, such as ADH and DCIS (Brancati et al., 2022). These results make the gap plausible, though their training budgets and evaluation units differ. This study computes patient-macro recall from patch predictions.

Imbalance damage differs strongly between the two tasks (Figure 2a). On TCGA-UT it is small and accelerating: 0.57 pp at $\rho = 10$, the 1-pp threshold only at $\rho = 20$, and 2.91 pp [2.48, 3.34] at $\rho = 100$. On BRACS it starts immediately and grows steadily by about 1.15 pp per doubling of ρ : 1.15 pp already at $\rho = 2$, 4.02 pp at $\rho = 10$, and 7.65 pp [5.77, 9.78] at $\rho = 100$, a fifth of the balanced level. Both native allocations lie on their curves (1.18 and 4.05 pp). The aggregate hides a large shift of recall from tail to head: at $\rho = 100$, the TCGA-UT tail third loses 11.8 pp while the head gains 5.4 pp; on BRACS the two tail classes drop from 38.7% to 7.3% recall. The predicted probabilities suffer far more than accuracy (Figure 2b). The expected calibration error (ECE), the average gap between the model’s confidence and its accuracy, rises from 5.5 to 17.8 pp on TCGA-UT and from 12.3 to 35.7 pp on BRACS at $\rho = 100$: the model becomes overconfident, but uniformly so. A global confidence correction that divides all logits by one scalar fitted on validation data (temperature scaling; Guo et al., 2017) removes 92–96% of the increase.

4 Class influence and tail support

When a tail class receives fewer training patches, two things change. Its errors count less often in the training loss, so common classes have more influence on the fitted decision boundary. This is the *frequency* effect: it concerns how strongly the available examples of each class shape the classifier. The model also sees fewer distinct tail patches, at $\rho = 100$ only one or two per patient, so it has less evidence about how that class varies. This is the *support* effect: giving those few patches more weight cannot supply the missing examples. The conditions in Table 1 separate these effects. F keeps all patches from the balanced condition but gives head classes more weight, isolating the frequency effect. S keeps the sparse patches from the imbalanced condition but weights classes equally, isolating the support effect. The interaction $I = \Delta_r - \Delta_F - \Delta_S$ measures any additional damage when both occur together.

The two tasks break down differently (Figure 3). On TCGA-UT the frequency effect carries the damage at moderate ratios (0.73 of 1.01 pp at $\rho = 20$) while thin support alone costs nothing; at $\rho = 100$ the two compound: frequency 1.20 pp [0.87, 1.53], support 0.53 pp [0.36, 0.70], and a clearly positive interaction of 1.18 pp [1.00, 1.37]. On BRACS both factors are material at $\rho = 100$, frequency 4.47 pp [2.84, 6.30] and support 3.29 pp [1.99, 4.67], and their sum matches the total damage. The interaction

Class influence	Many tail patches	Few tail patches
Equal	balanced (ref.)	S: support only
Head-favoured	F: frequency only	r: ratio (both)

Table 1: Four training conditions separate the effect of unequal class influence (frequency) from the effect of fewer tail examples (support) at each ρ .




Figure 3: Damage at $\rho = 100$ by component; diamonds give the total with its 95% interval.

is near zero, -0.11 pp, but its interval $[-1.74, 1.36]$ cannot rule out compounding of up to 1.4 pp. Thin support alone thus costs BRACS more than the entire imbalance damage on TCGA-UT.

5 Dependence on tail-class identity

On BRACS, rotating every class through every rank shows that easy classes cost most when rare: damage at $\rho = 100$ ranges from 5.73 pp (UDH in the tail) to 9.99 pp (invasive carcinoma), because a class with high balanced recall has the most to lose (Figure 4a). On TCGA-UT, a class is especially vulnerable when it receives few patches but a class it is often mistaken for receives many. A score fixed before testing new class orders used this pattern to predict their damage. The order ranked most damaging costs 4.51 pp, the least damaging 1.34 pp, and random orders 2.91 pp (Figure 4b). Reversing the most damaging order keeps the same rank gaps between pairs of classes but lowers damage by 1.11 pp $[0.47, 1.73]$. Damage is therefore greatest when the class prone to being mistaken for another is the rarer one.




Figure 4: Damage at $\rho = 100$ with 95% intervals. (a) BRACS with each class in the tail in turn, against its balanced-condition recall (Spearman 0.96 , $n = 7$). (b) TCGA-UT class-to-rank assignments.

6 Encoder comparison

Refitting all conditions on identical patients with UNI2-h (Chen et al., 2024) instead of Virchow2 leaves the picture intact. On TCGA-UT, UNI2-h starts 1.78 pp higher and loses 0.86 pp less at $\rho = 100$ (97.5% interval 0.40 to 1.33), a resolved but sub-1-pp difference. On BRACS the difference is unresolved and the tail classes collapse to about 9% recall with both encoders.

7 Implications for mitigation

- Rebalancing the training objective removes the frequency damage but cannot restore thin support. On BRACS, 3.3 pp remain at $\rho = 100$; recovering them needs more tail patches or patients, or methods that improve tail-class estimation.
- Confidence correction removes most calibration damage; mitigation must beat that baseline.
- At the same ρ , damage varies 2–3-fold with dataset and class order; randomize class order.

A Guide to the supporting experiments

The numbered folders below are under `experiments/03_class_imbalance/`. Each folder's `report/` directory contains a PDF and its LaTeX source. The reports give the protocols, full tables, and uncertainty estimates behind the results condensed here.

Headline results.

- **Imbalance curves and probability quality:** Reports 25 and 26, in `00_damage/25_damage_tcga_ut/report/` and `00_damage/26_damage_bracs/report/`, give the balanced baselines, damage at every ratio, native allocations, head-to-tail recall, and calibration before and after temperature scaling.
- **Class influence versus tail support:** Reports 27 and 28, in `01_cause/tcga_ut/27_prior_vs_support/report/` and `01_cause/bracs/28_prior_vs_support/report/`, show the paired frequency-only (called prior-only, P, in the reports) and support-only conditions, their interaction, and its uncertainty on each dataset.
- **Which classes become rare:** Report 29, in `01_cause/bracs/29_tail_assignment/report/`, rotates every BRACS class through the tail. Report 32, in `01_cause/tcga_ut/32_worst_assignment_rate/report/`, tests predicted damaging and mild class orders on TCGA-UT, including reversal of the most damaging order.
- **Encoder comparison:** Report 39, in `01_cause/cross/39_encoder_transfer/report/`, compares Virchow2 and UNI2-h on matched patients and patches and examines sensitivity to regularization selection.

Further reading. Reports 25 and 26 also track negative log-likelihood and fitted confidence corrections across the full imbalance grid. Reports 29 and 32 give class-level recall and confusion analyses behind the aggregate class-order effects, useful when the identity of a rare class matters. Report 39 adds per-class encoder comparisons and a fixed-regularization analysis, showing how the probe-selection rule affects the encoder contrast.

References

- N. Brancati, A. M. Anniciello, P. Pati, D. Riccio, G. Scognamiglio, G. Jaume, G. De Pietro, M. Di Bonito, A. Foncubierta, G. Botti, M. Gabrani, F. Feroce, and M. Frucci. BRACS: A Dataset for BReAst Carcinoma Subtyping in H&E Histology Images. *Database*, 2022:baac093, 2022.
- M. Buda, A. Maki, and M. A. Mazurowski. A systematic study of the class imbalance problem in convolutional neural networks. *Neural Networks*, 106:249–259, 2018.
- R. J. Chen, T. Ding, M. Y. Lu, et al. Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30(3):850–862, Mar. 2024.
- C. Cong, Y. Yang, S. Liu, et al. Imbalanced Histopathology Image Classification Using Deep Feature Graph Attention Network. In *2022 IEEE 19th International Symposium on Biomedical Imaging (ISBI)*, pages 1–4, 2022.
- C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger. On Calibration of Modern Neural Networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330, 2017.
- L. Ju, S. Yan, Y. Zhou, et al. MONICA: Benchmarking on Long-tailed Medical Image Classification. *arXiv (Cornell University)*, 2024.
- D. Komura, A. Kawabe, K. Fukuta, et al. Universal encoding of pan-cancer histology by deep texture representations. *Cell Reports*, 38(9):110424, 2022.
- Mahmood Lab. UNI: Towards a General-Purpose Foundation Model for Computational Pathology – UNI 2 ROI benchmarks. <https://github.com/mahmoodlab/UNI>, 2025.
- C. Mosquera, L. Ferrer, D. H. Milone, et al. Class imbalance on medical image classification: towards better evaluation practices for discrimination and calibration performance. *European Radiology*, 34(12):7895–7903, 2024.
- A. Vaidya, A. Zhang, G. Jaume, A. H. Song, T. Ding, S. J. Wagner, M. Y. Lu, et al. Molecular-driven Foundation Model for Oncologic Pathology. *arXiv preprint arXiv:2501.16652*, 2025. doi: 10.48550/arXiv.2501.16652.
- E. Zimmermann, E. Vorontsov, J. Viret, et al. Virchow2: Scaling Self-Supervised Mixed Magnification Models in Pathology. *arXiv preprint arXiv:2408.13910*, 2024.