

Extending Class-Imbalance Mitigation Benchmarks to External Datasets: A Native BRACS Benchmark

Yannik Queisler

July 8, 2026

Abstract

A prior study benchmarked class-imbalance mitigation methods on TCGA-UT, a pan-cancer whole-slide-image cohort with a native, moderately imbalanced slide distribution (Queisler, 2026). This report extends the same benchmark machinery to additional datasets, keeping the frozen Virchow2 representation, the bounded evidence budget, and the patient-disjoint evaluation protocol fixed, and asks whether the method comparison remains meaningful when the data regime changes. The first external dataset is BRACS, an organ-specific H&E breast-lesion cohort evaluated on its native seven-subtype distribution; further datasets are planned. Both patch-feature and whole-slide-image (WSI) feature-bag method families are compared, each within its own prediction regime, and a plain cross-entropy no-mitigation baseline (CE, MIL CE) is added so mitigation gains are measured against an unweighted reference. On BRACS’s native seven-subtype distribution (head-to-tail $\approx 2.3:1$), CFAL leads the patch regime by a clear margin over both the unweighted CE baseline and the other mitigation methods, while ProGAN synthetic augmentation is the weakest patch method overall, falling slightly below unweighted CE. In the WSI-bag regime the unweighted MIL CE baseline is competitive with, and in several cases stronger than, dedicated mitigation methods—only focal loss and SC-MIL clearly exceed it. Both regimes remain substantially miscalibrated after temperature scaling relative to what is achievable on the pan-cancer TCGA-UT cohort, reflecting the smaller training pool and the difficulty of fine-grained breast-lesion subtyping.

Contents

1	Introduction	2
2	Datasets	2
2.1	BRACS	2
3	Methods	4
4	Evaluation Protocol	4
5	Results	5
6	Discussion	7
7	Conclusion	7

1 Introduction

A companion study compared class-imbalance mitigation methods under the native, moderately imbalanced slide distribution of TCGA-UT (head-to-tail $\approx 28:1$) across 32 cancer types (Queisler, 2026; Tomczak et al., 2015; Komura et al., 2022). That benchmark fixes a frozen Virchow2 encoder, a bounded per-slide evidence budget, and patient-disjoint splits, so that the only thing being compared is how each mitigation method reshapes minibatch exposure or the per-example loss. The natural next question is external validity: does the same benchmark machinery stay meaningful when it is pointed at a different dataset, with a different organ, a different label set, and a different native class distribution?

This report takes the first step by extending the benchmark to *native* external datasets rather than to a controlled reconstruction of TCGA-UT. The design is deliberately conservative: the representation family (frozen Virchow2), the 30-instance evidence budget, the patient-disjoint split protocol, the method families, and the validation-based model-selection rule are all carried over unchanged, and only the dataset is swapped. A dataset therefore acts as an external check on the benchmark rather than as a new experimental variable.

The first external dataset is BRACS, an H&E breast-pathology dataset with 547 WSIs and 4,539 regions of interest (ROIs) from 189 patients, annotated by consensus of three board-certified pathologists into three lesion types and seven subtypes (Brancati et al., 2022). The atypical categories are clinically important and comparatively scarce, so BRACS is a genuine class-imbalance setting without any resampling. Unlike TCGA-UT it is organ-specific, has far fewer labels, and provides lesion-focused ROI annotations; these differences are exactly what make it a useful external check. Additional datasets are planned and will slot into the same benchmark.

Both regimes from the companion study are carried over. Patch-feature methods include cross-entropy, class weighting, balanced sampling, focal loss, Scholz-style metric losses, CFAL, OKO set learning, and ProGAN synthetic augmentation (Scholz et al., 2024). WSI-bag methods include cross-entropy, class weighting, balanced sampling, focal loss, RankMix, SC-MIL, and MDE-MIL (Chen and Lu, 2023; Juyal et al., 2024; Ling et al., 2025). Patch and WSI-bag methods are reported separately because they consume different prediction units and answer regime-specific questions.

2 Datasets

The benchmark evaluates each dataset on its *native* class distribution: the observed per-class support is measured after a patient-disjoint preparation and used directly, without any resampling or power-law reconstruction. This report covers BRACS; further datasets will be added under the same protocol.

2.1 BRACS

BRACS is an organ-specific H&E breast-lesion dataset rather than a pan-cancer dataset, and its annotations are lesion-subtype labels rather than cancer-type labels. It contains 547 WSIs and 4,539 ROIs from 189 patients, annotated by consensus of three board-certified pathologists (Brancati et al., 2022). The labels form three lesion types—benign, atypical, and malignant—and seven subtypes: normal tissue (N), pathological benign (PB), usual ductal hyperplasia (UDH), flat epithelial atypia (FEA), atypical ductal hyperplasia (ADH), ductal carcinoma in situ (DCIS), and invasive carcinoma (IC).

The benchmark uses the seven-subtype target because it is the level at which class imbalance is most relevant. The atypical categories FEA and ADH are clinically important precancerous lesions and are less common than many benign or malignant samples, so collapsing BRACS to

the three lesion types would hide part of the rare-class problem. Patient identifiers from the BRACS summary sheet define the split unit, and the prepared benchmark enforces patient-disjoint train, validation, and test partitions across three seeds—matching the leakage-control standard used for TCGA-UT.

BRACS evidence is derived from the ROI set. Each annotated ROI is tiled deterministically into 256×256 RGB patches, capped at 30 tiles per ROI to match the evidence budget used for the companion study. Patch methods consume one frozen Virchow2 feature per ROI tile (Vorontsov et al., 2024; Zimmermann et al., 2024), represented as the concatenation of the CLS token and the mean of the patch tokens (2,560-dimensional vectors). The WSI-bag pipeline consumes grouped ROI-tile features as bounded bags, preserving the same patch-versus-bag comparison while avoiding an additional tissue-detection step on full WSIs. This choice prioritizes label fidelity: all training instances originate from pathologist-selected lesion regions rather than from unlabeled background tissue.

After patch extraction, 387 of the 547 published WSIs contribute at least one annotated ROI tile; the remaining slides have no tiled ROI annotations available. The prepared benchmark covers 151 patients, the full 4,539 ROI metadata records, and 98,282 extracted patches (Table 1). The seven-subtype WSI support spans PB (196 slides) at the head to DCIS (87 slides) at the tail, a native head-to-tail ratio of approximately 2.25:1 (Figure 1). This is mild imbalance relative to the pan-cancer TCGA-UT cohort, but it reflects the true clinical distribution of breast-lesion subtypes and is used as-is.

Property	Value
Patients	151
WSIs with ROI tiles	387
ROI metadata rows	4,539
ROI tiles	98,282
Subtype labels	N, PB, UDH, FEA, ADH, DCIS, IC
WSI support range	87–196
WSI imbalance ratio	$\approx 2.3:1$

Table 1: BRACS dataset properties used by the native benchmark, after patient-disjoint preparation and ROI tiling.




Figure 1: Native BRACS WSI support by seven-subtype label after patient-disjoint preparation. The distribution is used directly as the benchmark’s native class-imbalance setting.

3 Methods

The method taxonomy follows the companion benchmark (Queisler, 2026). Cross-entropy is the no-mitigation reference. Data-level methods change minibatch exposure through balanced sampling or synthetic augmentation. Algorithm-level methods change the per-example loss through class weighting, focal loss, metric-derived objectives, prototype-affinity objectives, or OKO set supervision. WSI-specific methods change how slide bags are represented or mixed.

Family	Methods	Regime
No-mitigation baseline	CE, MIL CE	Patch, WSI bag
Data-level sampling and augmentation	Balanced sampling, RankMix	Patch, WSI bag
Algorithm-level losses	Weighted CE, Focal loss, CFAL	Patch (CFAL); Patch and WSI bag (others)
Hybrid sampling-loss	CE + soft F1, CE + soft MCC, OKO, MDE-MIL	Patch (first three, OKO); WSI bag (MDE-MIL)
Synthetic data generation	ProGAN augmentation	Patch
Representation and architecture	SC-MIL	WSI bag

Table 2: Method families evaluated on the native BRACS benchmark, following the taxonomy of Queisler (2026). CE and MIL CE are unweighted cross-entropy baselines, added so mitigation gains are measured against a genuine no-mitigation reference. Patch and WSI-bag methods are compared only within their own input regime.

Full method definitions follow Queisler (2026) and Scholz et al. (2024). RankMix, SC-MIL, and MDE-MIL are WSI-bag methods whose mechanisms operate on bags rather than isolated patch embeddings (Chen and Lu, 2023; Juyal et al., 2024; Ling et al., 2025). The plain cross-entropy baselines (CE for the patch regime, MIL CE for the WSI-bag regime) apply uniform class weights and no resampling, so any improvement from a mitigation method is measured relative to doing nothing.

4 Evaluation Protocol

The benchmark evaluates each dataset on its native class distribution. Hyperparameters are selected independently for every method: for each candidate configuration the validation macro F1 is averaged over the three patient-disjoint seeds, and the configuration with the highest mean is chosen, with ties broken by mean validation balanced accuracy. A single configuration is therefore fixed per method (shared across seeds). Selected configurations are evaluated on the fixed test split, and results are reported as mean $\pm$ standard deviation across seeds. Macro F1 is the primary discrimination metric because it weights classes equally; balanced accuracy and accuracy are reported as complementary discrimination metrics.

The comparison is intentionally regime-specific. Patch-feature methods are compared only with other patch-feature methods, and WSI-bag methods are compared only with other WSI-bag methods. The two regimes share the patient-disjoint splits, the frozen Virchow2 representation family, and the 30-instance evidence budget, but they differ in prediction unit and aggregation mechanism.

Calibration is evaluated alongside discrimination using expected calibration error (ECE; Naeini et al., 2015; Guo et al., 2017). For each method and seed, a single temperature is fit by minimizing negative log likelihood on the validation split and applied unchanged to the test score matrix before softmax (Guo et al., 2017). Tables 5 and 6 report ECE before and after temperature scaling with the fitted temperature.

5 Results

The patch-feature and WSI-bag regimes are reported separately (Tables 3 and 4). Scores are specific to the seven-class breast-lesion subtyping task and are not directly comparable to the 32-class TCGA-UT benchmark: the label set, task difficulty, and cohort size all differ.

Method	Accuracy	Balanced accuracy	Macro F1
CE	0.515 ± 0.070	0.457 ± 0.025	0.443 ± 0.035
Weighted CE	0.523 ± 0.054	0.468 ± 0.036	0.458 ± 0.038
Focal loss	0.504 ± 0.057	0.461 ± 0.032	0.446 ± 0.040
CFAL	0.581 ± 0.082	0.512 ± 0.050	0.499 ± 0.059
CE + soft F1	0.543 ± 0.059	0.490 ± 0.031	0.474 ± 0.035
CE + soft MCC	0.519 ± 0.048	0.467 ± 0.022	0.452 ± 0.024
Balanced sampling	0.513 ± 0.053	0.458 ± 0.035	0.447 ± 0.035
OKO	0.527 ± 0.064	0.482 ± 0.034	0.466 ± 0.036
ProGAN augmentation	0.498 ± 0.045	0.455 ± 0.021	0.437 ± 0.023

Table 3: BRACS native patch-feature test results (accuracy, balanced accuracy, and macro F1) for validation-selected configurations, mean $\pm$ standard deviation across patient-disjoint seeds.

Method	Accuracy	Balanced accuracy	Macro F1
MIL CE	0.603 ± 0.030	0.446 ± 0.043	0.430 ± 0.014
Weighted CE	0.577 ± 0.089	0.415 ± 0.049	0.423 ± 0.055
Focal loss	0.613 ± 0.070	0.469 ± 0.104	0.459 ± 0.060
Balanced sampling	0.617 ± 0.044	0.400 ± 0.061	0.395 ± 0.066
RankMix	0.579 ± 0.127	0.408 ± 0.084	0.414 ± 0.083
SC-MIL	0.633 ± 0.087	0.453 ± 0.063	0.472 ± 0.056
MDE-MIL	0.612 ± 0.087	0.426 ± 0.076	0.422 ± 0.064

Table 4: BRACS native WSI-bag test results for validation-selected configurations, mean $\pm$ standard deviation across patient-disjoint seeds.

In the patch regime, CFAL leads on all three metrics (BAcc 0.512, F1 0.499), 2.2 points of balanced accuracy ahead of the next-best method (CE + soft F1, BAcc 0.490) and 5.5 points ahead of the unweighted CE baseline (BAcc 0.457, F1 0.443). The remaining methods—Weighted CE, focal loss, CE + soft MCC, and balanced sampling—cluster tightly between BAcc 0.458 and 0.468, indistinguishable from the CE baseline within one standard deviation (all $\sigma \approx 0.03$ –0.08). OKO (BAcc 0.482) sits between this cluster and CFAL. ProGAN augmentation is the weakest method in the regime (BAcc 0.455, F1 0.437), fractionally below the unweighted CE baseline on every metric—unlike on the companion TCGA-UT benchmark, where ProGAN led at the mildest severity, BRACS’s native imbalance never becomes severe enough for synthetic augmentation to pay off. The overall picture is that most patch mitigation methods provide little measurable lift over unweighted cross-entropy on BRACS’s mild native imbalance, and one (ProGAN) falls slightly below it; only CFAL and, to a lesser extent, CE + soft F1 and OKO separate from the no-mitigation baseline.

In the WSI-bag regime, the unweighted MIL CE baseline (BAcc 0.446, F1 0.430) is com-

petitive with most dedicated mitigation methods: it outperforms Weighted CE (BAcc 0.415), balanced sampling (BAcc 0.400), RankMix (BAcc 0.408), and MDE-MIL (BAcc 0.426) on balanced accuracy. Only focal loss (BAcc 0.469, the regime leader) and SC-MIL (BAcc 0.453, F1 0.472, the regime’s highest macro F1) clearly exceed the CE baseline. Between-seed standard deviations are large throughout this regime (up to $\sigma \approx 0.10$ – 0.13 for RankMix and SC-MIL), so individual rankings should be read as numerical leads rather than statistically separated outcomes; the qualitative pattern—that several bag-composition and resampling methods do not clearly beat plain MIL CE on BRACS—is nonetheless the direct motivation for including the baseline. Balanced sampling is the weakest method in this regime (BAcc 0.400, F1 0.395), below the unweighted baseline on every metric.

Method	ECE	ECE+TS	T
CE	0.376 ± 0.058	0.056 ± 0.036	6.361 ± 1.353
Weighted CE	0.372 ± 0.044	0.047 ± 0.025	6.009 ± 1.212
Focal loss	0.280 ± 0.041	0.038 ± 0.029	3.300 ± 0.672
CFAL	0.418 ± 0.085	0.068 ± 0.032	0.058 ± 0.010
CE + soft F1	0.416 ± 0.054	0.078 ± 0.010	9.079 ± 0.906
CE + soft MCC	0.397 ± 0.045	0.058 ± 0.028	7.003 ± 1.266
Balanced sampling	0.378 ± 0.046	0.046 ± 0.028	6.069 ± 1.164
OKO	0.146 ± 0.055	0.043 ± 0.035	1.468 ± 0.312
ProGAN augmentation	0.394 ± 0.042	0.053 ± 0.031	6.337 ± 1.146

Table 5: BRACS native patch-feature calibration: ECE before (raw) and after (TS) temperature scaling, with the fitted temperature T , mean $\pm$ standard deviation across seeds.

Method	ECE	ECE+TS	T
MIL CE	0.329 ± 0.063	0.155 ± 0.002	3.301 ± 0.274
Weighted CE	0.332 ± 0.086	0.155 ± 0.034	3.313 ± 0.428
Focal loss	0.283 ± 0.073	0.134 ± 0.027	2.838 ± 0.395
Balanced sampling	0.292 ± 0.063	0.142 ± 0.030	3.156 ± 0.248
RankMix	0.152 ± 0.077	0.133 ± 0.054	1.160 ± 0.060
SC-MIL	0.259 ± 0.101	0.112 ± 0.033	3.126 ± 0.312
MDE-MIL	0.275 ± 0.043	0.152 ± 0.047	1.986 ± 0.242

Table 6: BRACS native WSI-bag calibration: ECE before and after temperature scaling, with the fitted temperature T . Columns as in Table 5.

Probability Calibration. In the patch regime, raw (pre-scaling) miscalibration is severe for most methods: ECE ranges from 0.28 to 0.42, requiring large temperature corrections ($T \approx 3$ – 9) to reach post-scaling ECE of 0.04–0.08. OKO is the exception, with by far the lowest raw ECE (0.146) and the smallest fitted temperature ($T \approx 1.47$), consistent with its behaviour on the companion TCGA-UT benchmark. CFAL again shows the CFAL-specific pattern of a near-degenerate low temperature ($T \approx 0.058$) alongside the highest post-scaling ECE (0.068) of any patch method, echoing its affinity-based outputs being nearly flat prior to scaling. Focal loss achieves the lowest post-scaling ECE (0.038). ProGAN augmentation calibrates similarly to the unweighted baselines (raw ECE 0.394, $T \approx 6.34$, post-scaling ECE 0.053), showing no distinctive miscalibration pattern despite its weaker discrimination. In the WSI-bag regime, calibration is uniformly worse than the patch regime and than the companion TCGA-UT benchmark: post-scaling ECE ranges from 0.112 (SC-MIL, the best-calibrated method) to 0.155 (MIL CE

and Weighted CE, tied for worst). RankMix has the lowest raw ECE (0.152) and smallest temperature ($T \approx 1.16$) of any WSI-bag method, but its post-scaling ECE (0.133) is not the best. The elevated residual miscalibration across all BRACS methods, relative to TCGA-UT, reflects the smaller training pool and the intrinsic difficulty of seven-class fine-grained subtyping rather than a failure specific to any one method.

6 Discussion

The benchmark’s central design change from the companion TCGA-UT study is the addition of an unweighted no-mitigation baseline (CE, MIL CE), and the two regimes answer the question it was added to ask quite differently. In the patch regime, most mitigation methods do not clearly separate from the CE baseline on BRACS’s mild native imbalance ($\approx 2.3:1$)—Weighted CE, focal loss, CE + soft MCC, and balanced sampling all sit within noise of unweighted cross-entropy, and ProGAN augmentation falls fractionally below it. Only CFAL, and to a lesser extent CE + soft F1 and OKO, produce a measurable gain. This is a plausible consequence of the imbalance itself being mild: with a native ratio an order of magnitude below TCGA-UT’s $\approx 28:1$, there may simply be less headroom for reweighting or resampling to help. ProGAN’s reversal is particularly informative: on TCGA-UT it led at the mildest constructed severity before collapsing as tail-class support fell to single-digit patch counts at the harshest severity. BRACS’s native imbalance never reaches that regime, so ProGAN never gets the tail-scarcity conditions under which synthetic augmentation would be expected to help, and it is left carrying only the risk of the generative step (distributional mismatch between synthetic and real patches) without the corresponding benefit.

The WSI-bag regime tells a sharper version of the same story: several mitigation methods (Weighted CE, balanced sampling, RankMix, MDE-MIL) score *below* the unweighted MIL CE baseline on balanced accuracy. Only focal loss and SC-MIL clearly improve on it. This is exactly the failure mode a no-mitigation baseline is meant to surface, and it would have been invisible in a benchmark that only compared mitigation methods against each other—the companion TCGA-UT report, which lacked this baseline, could not have detected it. Whether this reflects BRACS’s smaller training pool (108 native WSIs after preparation, versus TCGA-UT’s constructed pools of over 1000 slides), the milder native imbalance giving resampling less to correct, or a genuine difference in how bag-composition methods behave outside a pan-cancer setting cannot be distinguished from a single external dataset; it is a hypothesis for the next dataset in this benchmark to test.

Calibration adds a second, independent axis on which BRACS is harder than TCGA-UT: even the best-calibrated method in each regime (focal loss for patch, SC-MIL for WSI-bag) needs temperature scaling to reach an ECE that is still $2\text{--}3\times$ worse than typical TCGA-UT post-scaling values. This is consistent with the smaller training pool and finer-grained label set rather than a method-specific weakness, since the pattern holds across every method in both regimes.

7 Conclusion

This report extended the class-imbalance mitigation benchmark introduced for TCGA-UT (Queisler, 2026) to a first external dataset, BRACS, keeping the frozen Virchow2 representation, the 30-instance evidence budget, and the patient-disjoint evaluation protocol unchanged and adding an unweighted cross-entropy baseline (CE, MIL CE) so mitigation gains are measured against doing nothing. Evaluated on BRACS’s native seven-subtype distribution ($\approx 2.3:1$ head-to-tail), the patch regime shows CFAL leading by a clear margin (BAcc 0.512 vs. 0.457 for CE), with most other mitigation methods statistically indistinguishable from the no-mitigation baseline and ProGAN augmentation—the leader at mild severity on the companion TCGA-UT

benchmark—falling fractionally below it (BAcc 0.455). The WSI-bag regime shows that several mitigation methods—Weighted CE, balanced sampling, RankMix, and MDE-MIL—do not clearly improve on the unweighted MIL CE baseline; only focal loss and SC-MIL do. Both regimes remain substantially more miscalibrated after temperature scaling than the companion TCGA-UT benchmark, consistent with BRACS’s smaller training pool and harder seven-class fine-grained subtyping task.

These results validate the benchmark’s external applicability while also showing that method rankings are not universal: CFAL’s patch-level lead, ProGAN’s reversal from TCGA-UT leader to BRACS laggard, and several WSI-bag methods’ failure to beat plain cross-entropy are all new findings relative to TCGA-UT, and none would have been visible without a native external dataset and a genuine no-mitigation reference. Further external datasets remain to be added under the same protocol; the controlled power-law severity construction previously explored for TCGA-UT and BRACS has been moved to a separate follow-up experiment (see `controlled_imbalance_experiment.md`) rather than folded into this native-benchmark report.

References

- Nadia Brancati, Anna Maria Anniciello, Pushpak Pati, Daniel Riccio, Giosue Scognamiglio, Guillaume Jaume, Giuseppe De Pietro, Maurizio Di Bonito, Antonio Foncubierta, Gerardo Botti, Maria Gabrani, Florinda Feroce, and Maria Frucci. BRACS: A dataset for BReAst carcinoma subtyping in H&E histology images. *Database*, 2022:baac093, 2022. doi: 10.1093/database/baac093.
- Yuan-Chih Chen and Chun-Shien Lu. RankMix: Data augmentation for weakly supervised learning of classifying whole slide images with diverse sizes and imbalanced categories. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23936–23945, 2023. doi: 10.1109/CVPR52729.2023.02397.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330. PMLR, 2017.
- Dinkar Juyal, Siddhant Shingi, Syed Ashar Javed, Harshith Padigela, Chintan Shah, Anand Sampat, Archit Khosla, John Abel, and Amaro Taylor-Weiner. SC-MIL: Supervised contrastive multiple instance learning for imbalanced classification in pathology. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 7016–7025, 2024. doi: 10.1109/WACV57761.2024.00702.
- Daisuke Komura, Akihiro Kawabe, Kenta Fukuta, Kazuya Sano, Tsukasa Umezaki, Haruka Koda, Riku Suzuki, Keisuke Tominaga, Michiaki Ochi, Hirofumi Konishi, et al. Universal encoding of pan-cancer histology by deep texture representations. *Cell Reports*, 38(9):110424, 2022. doi: 10.1016/j.celrep.2022.110424.
- Xitong Ling, Yifeng Ping, Jiawen Li, Jing Peng, Yuxuan Chen, Minxi Ouyang, Yizhi Wang, Yonghong He, Tian Guan, Xiaoping Liu, and Lianghui Zhu. Multimodal distillation-driven ensemble learning for long-tailed histopathology whole slide images analysis. *arXiv preprint arXiv:2503.00915*, 2025. doi: 10.48550/arXiv.2503.00915.
- Mahdi Pakdaman Naeini, Gregory Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using bayesian binning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 29, pages 2901–2907, 2015. doi: 10.1609/aaai.v29i1.9602.
- Yannik Queisler. Class imbalance mitigation for TCGA-UT. Technical report, TU Berlin, 2026.
- Daniel Scholz, Ayhan Can Erdur, Josef Buchner, Jan C. Peeken, Daniel Rueckert, and Benedikt Wiestler. Imbalance-aware loss functions improve medical image classification. In *Proceedings of the 7th International Conference on Medical Imaging with Deep Learning*, volume 250 of *Proceedings of Machine Learning Research*, pages 1341–1356. PMLR, 2024.
- Katarzyna Tomczak, Patrycja Czerwinska, and Maciej Wiznerowicz. The cancer genome atlas (TCGA): an immeasurable source of knowledge. *Contemporary Oncology*, 19(1A):A68–A77, 2015. doi: 10.5114/wo.2014.47136.
- Eugene Vorontsov, Alican Bozkurt, Adam Casson, George Shaikovski, Michal Zelechowski, Kristen Severson, Eric Zimmermann, et al. A foundation model for clinical-grade computational pathology and rare cancers detection. *Nature Medicine*, 30(10):2924–2935, 2024. doi: 10.1038/s41591-024-03141-0.
- Eric Zimmermann, Eugene Vorontsov, Julian Viret, Adam Casson, Michal Zelechowski, George Shaikovski, Neil Tenenholtz, James Hall, Thomas Fuchs, Nicolo Fusi, Siqi Liu, and Kristen Severson. Virchow2: Scaling self-supervised mixed magnification models in pathology. *arXiv preprint arXiv:2408.00738*, 2024. doi: 10.48550/arXiv.2408.00738.