

Controlled Class-Imbalance Construction for TCGA-UT

Yannik Queisler

June 23, 2026

Abstract

Class imbalance in computational pathology is routinely studied under real distributions, where prevalence, class separability, and data-collection history are entangled. This report defines a controlled TCGA-UT benchmark that varies only the imbalance severity of the training distribution while holding validation and test splits fixed. Power-law training distributions are generated under native-prevalence class ordering at imbalance parameters $\lambda \in \{0.5, 1.0, 1.5\}$ with a shared strictly feasible pool of 1,132 slides, yielding head-to-tail ratios of approximately 6:1, 30:1, and 122:1 that bracket and extend the native TCGA-UT imbalance ratio ($\approx 28:1$). Using frozen Virchow2 patch and WSI-bag features with a 30-instance evidence budget, the same method families evaluated in a prior class-imbalance study on TCGA-UT (Queisler, 2026) are compared across three constructed severities. Method rankings are broadly preserved at patch level, while balanced accuracy declines by eight to nine percentage points from $\lambda = 0.5$ to $\lambda = 1.5$; WSI-bag methods show larger absolute performance and RankMix is the most severity-robust bag-composition method.

Contents

1	Introduction	2
2	Dataset Construction	2
2.1	Power-Law Training Support	3
2.2	Feasible Pool Size and Realized Profiles	3
3	Methods	3
4	Evaluation Protocol	4
5	Results	5
6	Discussion	7
7	Conclusion	7
A	Full Calibration Metrics	10

1 Introduction

A prior class-imbalance study on TCGA-UT compared mitigation methods under the dataset’s native, moderately imbalanced slide distribution (head-to-tail $\approx 28:1$) across 32 cancer types (Queisler, 2026). Native distributions are realistic but confound several factors simultaneously: how often each class appears, tumor morphology, tissue diversity, slide acquisition, and the representation geometry of the frozen encoder. A controlled construction isolates one factor—training-support frequency—by imposing known power-law distributions while holding validation and test data fixed.

TCGA-UT is a suitable test bed for this isolation because it provides slide labels and patch-level evidence from diagnostic whole-slide images in The Cancer Genome Atlas (Tomczak et al., 2015; Komura et al., 2022). Controlling only the imbalance severity under a fixed native-prevalence class ordering allows the question to be posed cleanly: do mitigation methods maintain their relative advantage as the imbalance between frequent and rare cancer types intensifies along the same ranking?

The report evaluates the same method families as Queisler (2026) under three constructed training distributions that vary only $\lambda \in \{0.5, 1.0, 1.5\}$, the imbalance-severity parameter, while holding the class ranking fixed. Patch-level methods include cross-entropy, class weighting, balanced sampling, focal loss, Scholz-style metric losses, CFAL, and OKO set learning (Scholz et al., 2024). WSI-bag methods include cross-entropy, class weighting, balanced sampling, focal loss, RankMix, SC-MIL, and MDE-MIL (Chen and Lu, 2023; Juyal et al., 2024; Ling et al., 2025). Patch and WSI-bag methods are reported separately because they consume different prediction units and answer regime-specific questions.

2 Dataset Construction

The construction shares the TCGA-UT case-disjoint split protocol with Queisler (2026). For each of three random seeds, the validation and test splits are held fixed, and artificial imbalance is applied only to the training split. The assignment is patient-disjoint: all slides and feature rows associated with one TCGA participant remain in exactly one partition. This isolates the effect of training-support frequency: model selection and evaluation are comparable across imbalance parameters and mitigation methods without patient leakage.

All 32 cancer types are retained. Let $C = 32$, let $\mathcal{D}_{\text{train}}$ denote the seed-specific training slide pool, and let a_c be the number of available training slides for class c . The constructed training split must satisfy

$$1 \leq n_c \leq a_c, \quad \sum_{c=1}^C n_c = N, \quad c \in \{1, \dots, C\}. \quad (1)$$

Every class receives at least one slide, no class is oversampled beyond its available pool, and the total training set is fixed at N slides shared across all imbalance parameters and seeds.

After slide selection, the input evidence is capped. Patch classifiers use up to 30 deterministic patches per slide. WSI-bag classifiers use up to 30 deterministic frozen feature instances per slide. In both regimes, Virchow2 features remain frozen (Vorontsov et al., 2024; Zimmermann et al., 2024), read from the shared `cls_patchmean` store (`concat(CLS, mean(patch tokens))`, 2,560-dimensional vectors). The cap aligns the input budget across regimes and prevents slides with more extracted tumor patches from dominating the training objective.

2.1 Power-Law Training Support

For an ordered class list $(c_1, \dots, c_C)$, the intended share for rank i under imbalance parameter λ is

$$p_i(\lambda) = \frac{i^{-\lambda}}{\sum_{j=1}^C j^{-\lambda}}. \quad (2)$$

When $\lambda = 0$ the intended distribution is uniform over ranks; as λ increases, early-ranked classes receive more training support and late-ranked classes receive less. All constructed splits use the *native-prevalence order*: classes are ranked by their TCGA-UT training-split slide count, so rank 1 is the most frequent cancer type and rank $C = 32$ the least frequent. The power-law form is imposed to obtain controlled, monotonic severity steps; it is not a fit to the native TCGA-UT distribution, which is only moderately imbalanced ($\approx 28:1$) and not well described by a single power law.

Raw targets $N p_i(\lambda)$ are converted to integer counts by flooring and assigning residual slides by largest fractional remainder. Each class first receives one guaranteed slide; the remaining $N - C$ slides are then allocated according to Equation (2). A regime is accepted only when every rounded target satisfies Equation (1).

2.2 Feasible Pool Size and Realized Profiles

The pool size N must be small enough that no class target exceeds its available support across all evaluated imbalance parameters. Define the maximum feasible pool for a given λ and seed as

$$T^*(\lambda) = \min_c \frac{a_c}{p_c(\lambda)}. \quad (3)$$

For $N \leq T^*(\lambda)$, all targets $N p_c(\lambda)$ are feasible. Setting $N = \lfloor \min_s T_s^*(0.5) \rfloor$ over the three seeds guarantees exact feasibility at the mildest evaluated imbalance on every seed. The binding regime is $\lambda = 0.5$ rather than the steepest: the flat power-law profile at $\lambda = 0.5$ requires every class—including rare ones with limited available slides—to contribute a non-trivial share, saturating the rarest class’s supply at a lower total N than the peaked profiles at $\lambda \geq 1.0$. For the TCGA-UT training splits, this yields $N = 1,132$ slides, approximately 18% of the full available training pool.

The three λ values produce head-to-tail ratios of approximately 6:1 ($\lambda = 0.5$), 30:1 ($\lambda = 1.0$), and 122:1 ($\lambda = 1.5$) (Figure 1). The $\lambda = 0.5$ setting places the constructed distribution well below the native TCGA-UT ratio of approximately 28:1; $\lambda = 1.0$ reproduces near-native severity; and $\lambda = 1.5$ is substantially more skewed than anything observed in the native data.

3 Methods

The method taxonomy follows the companion benchmark, but this report uses it to study controlled training distributions rather than only the native TCGA-UT distribution. Cross-entropy is the no-mitigation reference. Data-level methods change minibatch exposure through balanced sampling or synthetic augmentation. Algorithm-level methods change the per-example loss through class weighting, focal loss, metric-derived objectives, prototype-affinity objectives, or OKO set supervision. WSI-specific methods change how slide bags are represented or mixed.

Full method definitions follow Queisler (2026) and Scholz et al. (2024). RankMix, SC-MIL, and MDE-MIL are WSI-bag methods whose mechanisms operate on bags rather than isolated patch embeddings (Chen and Lu, 2023; Juyal et al., 2024; Ling et al., 2025).




Figure 1: Constructed training support by class rank under the three evaluated imbalance parameters, with the native full-pool distribution shown as a reference. Classes are ordered by native prevalence (rank 1 = most frequent).

Family	Methods	Regime
No-mitigation baseline	CE / MIL CE	Patch, WSI bag
Data-level sampling and augmentation	Balanced sampling, RankMix	Patch, WSI bag
Algorithm-level losses	Weighted CE, Focal loss, CFAL	Patch (CFAL); Patch and WSI bag (others)
Hybrid sampling-loss	CE + soft F1, CE + soft MCC, OKO, MDE-MIL	Patch (first three, OKO); WSI bag (MDE-MIL)
Synthetic data generation	ProGAN augmentation	Patch
Representation and architecture	SC-MIL	WSI bag

Table 1: Method families evaluated under constructed TCGA-UT training distributions, following the taxonomy of Queisler (2026). Patch and WSI-bag methods are compared only within their own input regime.

4 Evaluation Protocol

Each constructed split is indexed by seed and imbalance parameter $\lambda \in \{0.5, 1.0, 1.5\}$. For every method and constructed regime, hyperparameters are selected by validation macro F1 across three seeds; ties are broken by validation balanced accuracy. Selected configurations are evaluated on the fixed test split, and results are reported as mean $\pm$ standard deviation across seeds. Macro F1 is the primary discrimination metric because it weights cancer types equally. Balanced accuracy and accuracy are reported as complementary discrimination metrics.

The comparison is intentionally regime-specific. Patch-feature methods are compared only with other patch-feature methods, and WSI-bag methods are compared only with other WSI-bag methods. The two regimes share the case-disjoint splits, frozen Virchow2 representation family, and 30-instance evidence budget, but they differ in prediction unit and aggregation mechanism.

Calibration is evaluated alongside discrimination using negative log likelihood (NLL), multiclass Brier score (Brier, 1950), and expected calibration error (ECE; Naeini et al., 2015; Guo

et al., 2017). For each method and seed, a single temperature is fit by minimizing NLL on the validation split and applied unchanged to the test score matrix before softmax (Guo et al., 2017). Tables 5 and 6 report ECE before and after temperature scaling; full NLL and Brier metrics are given in Tables 7 and 8.

5 Results

The analysis asks how methods respond as imbalance severity increases from $\lambda = 0.5$ to $\lambda = 1.5$ under the native-prevalence class order. The realized support profiles are summarised in Table 2; Tables 3 and 4 compare mitigation methods within the patch-feature and WSI-bag regimes.

λ	Training slides	Support range	Head:tail ratio
0.5	1,132	20–112	$\approx 6:1$
1.0	1,132	9–272	$\approx 30:1$
1.5	1,132	4–487	$\approx 122:1$
native (full pool)	6,110	20–567	$\approx 28:1$

Table 2: Constructed training-support profiles under native-prevalence ordering. Support range and head:tail ratio are the minimum and maximum per-class slide counts averaged across seeds. The native full-pool row is shown for reference; it is not part of the constructed experiment.

Method	$\lambda = 0.5$		$\lambda = 1.0$		$\lambda = 1.5$	
	BAcc	F1	BAcc	F1	BAcc	F1
Weighted CE	0.721 ± 0.012	0.693 ± 0.006	0.704 ± 0.012	0.693 ± 0.016	0.634 ± 0.006	0.641 ± 0.009
Focal loss	0.718 ± 0.011	0.692 ± 0.007	0.698 ± 0.016	0.689 ± 0.018	0.630 ± 0.005	0.641 ± 0.009
CFAL	0.724 ± 0.009	0.697 ± 0.007	0.710 ± 0.014	0.704 ± 0.015	0.632 ± 0.013	0.645 ± 0.014
CE + soft F1	0.721 ± 0.012	0.685 ± 0.004	0.713 ± 0.007	0.686 ± 0.014	0.638 ± 0.021	0.630 ± 0.019
CE + soft MCC	0.719 ± 0.006	0.690 ± 0.009	0.698 ± 0.010	0.684 ± 0.009	0.634 ± 0.019	0.633 ± 0.018
Balanced sampling	0.719 ± 0.012	0.691 ± 0.007	0.701 ± 0.015	0.688 ± 0.018	0.631 ± 0.012	0.635 ± 0.013
OKO	0.722 ± 0.012	0.689 ± 0.008	0.712 ± 0.011	0.694 ± 0.018	0.644 ± 0.015	0.648 ± 0.014

Table 3: Patch-feature test results for validation-selected configurations under constructed training distributions.

Patch-feature results show a consistent, monotone decline as imbalance intensifies (Table 3). At the mildest severity ($\lambda = 0.5$), CFAL leads in balanced accuracy (0.724), followed closely by OKO (0.722); at moderate severity ($\lambda = 1.0$), CE+soft-F1 (0.713) and OKO (0.712) move to the front; and at the harshest regime ($\lambda = 1.5$), OKO (0.644) retains a narrow lead over CE+soft-F1 (0.638), while focal loss is consistently the weakest patch method. Individual declines from $\lambda = 0.5$ to $\lambda = 1.5$ span eight to nine percentage points. Within the patch family the spread across methods at any fixed λ is approximately one percentage point—so the severity-induced decline is substantially larger than the advantage of choosing the best method.

The WSI-bag regime shows higher absolute performance than the patch regime (Table 4). MDE-MIL leads at mild and moderate severities (0.836 and 0.818), while RankMix overtakes it at the harshest regime (0.765 versus MDE-MIL’s 0.730). RankMix is the most severity-robust method overall, declining by approximately 6.5 percentage points from $\lambda = 0.5$ to $\lambda = 1.5$; focal loss is the least stable, declining by nearly 12 percentage points. At $\lambda = 1.5$, MDE-MIL falls narrowly below Weighted CE (0.736), indicating that the advantage of consistency regularisation over simple loss weighting erodes at the highest severity. The ordering of bag-composition methods (RankMix, SC-MIL) over loss-based variants is maintained throughout

Method	$\lambda = 0.5$		$\lambda = 1.0$		$\lambda = 1.5$	
	BAcc	F1	BAcc	F1	BAcc	F1
Weighted CE	0.823 ± 0.019	0.793 ± 0.017	0.799 ± 0.008	0.787 ± 0.007	0.736 ± 0.027	0.731 ± 0.021
Focal loss	0.821 ± 0.013	0.792 ± 0.012	0.801 ± 0.009	0.792 ± 0.007	0.703 ± 0.029	0.708 ± 0.031
Balanced sampling	0.809 ± 0.014	0.776 ± 0.010	0.797 ± 0.010	0.785 ± 0.018	0.724 ± 0.042	0.720 ± 0.040
RankMix	0.830 ± 0.016	0.793 ± 0.015	0.817 ± 0.024	0.795 ± 0.020	0.765 ± 0.029	0.751 ± 0.019
SC-MIL	0.813 ± 0.015	0.781 ± 0.014	0.800 ± 0.013	0.787 ± 0.008	0.745 ± 0.022	0.743 ± 0.015
MDE-MIL	0.836 ± 0.013	0.805 ± 0.012	0.825 ± 0.009	0.814 ± 0.007	0.730 ± 0.009	0.731 ± 0.014

Table 4: WSI-bag test results for validation-selected configurations under constructed training distributions.

all three regimes, confirming that WSI-specific aggregation mechanisms retain their advantage under increasing imbalance.

Method	$\lambda = 0.5$		$\lambda = 1.0$		$\lambda = 1.5$		T
	ECE	ECE+TS	ECE	ECE+TS	ECE	ECE+TS	
Weighted CE	0.156 ± 0.003	0.015 ± 0.002	0.157 ± 0.013	0.016 ± 0.005	0.174 ± 0.008	0.011 ± 0.006	2.43 ± 0.01
Focal loss	0.098 ± 0.003	0.010 ± 0.001	0.100 ± 0.012	0.012 ± 0.002	0.150 ± 0.008	0.013 ± 0.004	1.71 ± 0.22
CFAL	0.723 ± 0.009	0.067 ± 0.007	0.728 ± 0.021	0.070 ± 0.008	0.695 ± 0.013	0.077 ± 0.006	0.05 ± 0.00
CE + soft F1	0.179 ± 0.007	0.016 ± 0.004	0.184 ± 0.012	0.014 ± 0.002	0.217 ± 0.017	0.013 ± 0.003	3.22 ± 0.26
CE + soft MCC	0.162 ± 0.001	0.012 ± 0.004	0.169 ± 0.008	0.013 ± 0.002	0.183 ± 0.019	0.015 ± 0.005	2.51 ± 0.03
Balanced sampling	0.156 ± 0.005	0.016 ± 0.002	0.158 ± 0.013	0.014 ± 0.003	0.182 ± 0.011	0.017 ± 0.004	2.37 ± 0.05
OKO	0.071 ± 0.011	0.030 ± 0.003	0.193 ± 0.006	0.041 ± 0.004	0.192 ± 0.013	0.037 ± 0.007	0.70 ± 0.08

Table 5: Patch-feature calibration: ECE before (raw) and after (TS) temperature scaling across evaluated imbalance parameters. Fitted temperature T is averaged across seeds and severities. Full NLL and Brier metrics appear in Table 7.

Method	$\lambda = 0.5$		$\lambda = 1.0$		$\lambda = 1.5$		T
	ECE	ECE+TS	ECE	ECE+TS	ECE	ECE+TS	
Weighted CE	0.099 ± 0.012	0.022 ± 0.007	0.092 ± 0.010	0.022 ± 0.014	0.109 ± 0.009	0.023 ± 0.003	1.98 ± 0.05
Focal loss	0.049 ± 0.013	0.019 ± 0.004	0.054 ± 0.008	0.016 ± 0.005	0.105 ± 0.020	0.031 ± 0.004	1.46 ± 0.14
Balanced sampling	0.105 ± 0.010	0.025 ± 0.007	0.099 ± 0.012	0.020 ± 0.004	0.124 ± 0.020	0.027 ± 0.003	2.05 ± 0.06
RankMix	0.025 ± 0.015	0.020 ± 0.009	0.026 ± 0.005	0.018 ± 0.005	0.065 ± 0.008	0.026 ± 0.004	1.16 ± 0.15
SC-MIL	0.096 ± 0.006	0.026 ± 0.005	0.093 ± 0.011	0.018 ± 0.004	0.110 ± 0.001	0.024 ± 0.006	1.90 ± 0.04
MDE-MIL	0.051 ± 0.007	0.037 ± 0.005	0.042 ± 0.006	0.034 ± 0.007	0.044 ± 0.004	0.039 ± 0.010	0.96 ± 0.10

Table 6: WSI-bag calibration: ECE and fitted temperature T across evaluated imbalance parameters. Columns as in Table 5.

Probability Calibration. In the patch regime, CFAL is a calibration outlier: its affinity-based outputs are nearly flat over 32 classes, requiring strong temperature concentration ($T \approx 0.05$) and remaining substantially miscalibrated ($ECE \approx 0.07$) even after scaling—roughly five times higher than all other patch methods post-scaling. All other patch methods achieve ECE below 0.04 after temperature scaling. OKO shows a non-monotone calibration trajectory: well-calibrated at mild severity ($ECE \approx 0.07$ raw at $\lambda = 0.5$) but deteriorating to $ECE \approx 0.19$ at moderate severity before scaling. In the WSI-bag regime, RankMix is already well calibrated before scaling ($ECE \approx 0.025$); MDE-MIL shows only marginal improvement from temperature scaling and retains residual miscalibration ($ECE+TS \approx 0.037$ – 0.039). All WSI-bag methods reach ECE at or below 0.039 after scaling. Calibration rankings are broadly stable across imbalance severities.

6 Discussion

The controlled construction reveals two complementary patterns. First, method rankings are broadly preserved across imbalance severities, with moderate reshuffling among closely performing methods at the extremes. Second, balanced accuracy declines monotonically with λ for almost all methods, with a total range of eight to nine percentage points from the mildest to the strongest constructed regime. These patterns suggest that the relative utility of mitigation methods is largely robust to the particular severity chosen within this range, while absolute performance degrades substantially as rare-class support diminishes.

At the patch level, declines span from approximately 7.8 percentage points for OKO to 9.2 for CFAL. Within-method spread at any fixed λ is approximately one percentage point—so the severity-induced decline is substantially larger than the method-choice advantage. OKO and CE+soft-F1 maintain their relative advantage at the harshest severity; focal loss is consistently weakest throughout.

In the WSI-bag regime, RankMix stands out for its stability, declining by only 6.5 percentage points from $\lambda = 0.5$ to $\lambda = 1.5$. Focal MIL is the least stable, dropping by nearly 12 percentage points and landing at the bottom of the WSI rankings at $\lambda = 1.5$. MDE-MIL declines more than RankMix and falls narrowly below Weighted CE at the harshest regime, showing that consistency regularisation provides diminishing returns as rare-class support becomes critically scarce. The ordering of bag-composition methods (RankMix, SC-MIL) over simple loss-based variants holds throughout, confirming that WSI-specific aggregation mechanisms are severity-robust.

Calibration analysis complements discrimination: CFAL leads on patch discrimination but requires the strongest temperature correction and retains the highest post-scaling ECE. In the WSI-bag regime, RankMix is already well calibrated before scaling, while MDE-MIL shows only marginal improvement from temperature scaling despite leading on discrimination at mild and moderate severities.

Taken together, the controlled construction shows that when training-support frequency is imposed directly via a power-law distribution, aggregate discrimination declines predictably but modestly with severity, and the method ranking remains informative across the range tested. The prediction unit—patch versus WSI-bag—continues to be the dominant determinant of absolute performance: WSI-bag methods achieve balanced accuracy approximately ten percentage points higher than their patch-feature counterparts at equivalent imbalance.

7 Conclusion

This report defines a full-scale controlled TCGA-UT imbalance benchmark that keeps all 32 cancer types, preserves patient-disjoint validation and test splits, and varies only the training-support frequency under native-prevalence class ordering at $\lambda \in \{0.5, 1.0, 1.5\}$. The shared pool of 1,132 slides is determined by the mildest parameter $\lambda = 0.5$, whose flat profile saturates the rarest class’s available supply earliest. The three regimes span head-to-tail ratios of approximately 6:1, 30:1, and 122:1, bracketing and extending the native TCGA-UT ratio ($\approx 28:1$). Method rankings are broadly preserved across all three constructed severities, while balanced accuracy declines by eight to nine percentage points from the mildest to the strongest regime. At the patch level, CFAL and OKO lead at mild severity while OKO and CE+soft-F1 are most robust at the harshest; the within-method spread of approximately one percentage point at any fixed severity is substantially smaller than the severity-induced decline. At the WSI-bag level, bag-composition methods (RankMix, SC-MIL) consistently outperform loss-based variants regardless of severity; RankMix is the most severity-robust method, while MDE-MIL’s advantage over loss-based baselines erodes under the most extreme imbalance. The prediction unit remains the dominant performance determinant: WSI-bag methods achieve balanced accuracy

roughly 12 percentage points above patch-feature methods at equivalent imbalance. Temperature scaling is sufficient to align probability outputs across all methods, though CFAL requires the largest correction in the patch regime and MDE-MIL shows limited benefit from scaling in the WSI-bag regime.

References

- Glenn W. Brier. Verification of forecasts expressed in terms of probabilities. *Monthly Weather Review*, 78(1):1–3, 1950. doi: 10.1175/1520-0493(1950)078<0001:VOFEIT>2.0.CO;2.
- Yuan-Chih Chen and Chun-Shien Lu. RankMix: Data augmentation for weakly supervised learning of classifying whole slide images with diverse sizes and imbalanced categories. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23936–23945, 2023. doi: 10.1109/CVPR52729.2023.02397.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330. PMLR, 2017.
- Dinkar Juyal, Siddhant Shingi, Syed Ashar Javed, Harshith Padigela, Chintan Shah, Anand Sampat, Archit Khosla, John Abel, and Amaro Taylor-Weiner. SC-MIL: Supervised contrastive multiple instance learning for imbalanced classification in pathology. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 7016–7025, 2024. doi: 10.1109/WACV57761.2024.00702.
- Daisuke Komura, Akihiro Kawabe, Kenta Fukuta, Kazuya Sano, Tsukasa Umezaki, Haruka Koda, Riku Suzuki, Keisuke Tominaga, Michiaki Ochi, Hirofumi Konishi, et al. Universal encoding of pan-cancer histology by deep texture representations. *Cell Reports*, 38(9):110424, 2022. doi: 10.1016/j.celrep.2022.110424.
- Xitong Ling, Yifeng Ping, Jiawen Li, Jing Peng, and Yuxuan Chen. MDE-MIL. *Dos mil tres mil*, 27, 2025. doi: 10.35707/dostresmil/27.
- Mahdi Pakdaman Naeini, Gregory Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using bayesian binning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 29, pages 2901–2907, 2015. doi: 10.1609/aaai.v29i1.9602.
- Yannik Queisler. Class imbalance mitigation for TCGA-UT. Technical report, TU Berlin, 2026.
- Daniel Scholz, Ayhan Can Erdur, Josef Buchner, Jan C. Peeken, Daniel Rueckert, and Benedikt Wiestler. Imbalance-aware loss functions improve medical image classification. In *Proceedings of the 7th International Conference on Medical Imaging with Deep Learning*, volume 250 of *Proceedings of Machine Learning Research*, pages 1341–1356. PMLR, 2024.
- Katarzyna Tomczak, Patrycja Czerwinska, and Maciej Wiznerowicz. The cancer genome atlas (TCGA): an immeasurable source of knowledge. *Contemporary Oncology*, 19(1A):A68–A77, 2015. doi: 10.5114/wo.2014.47136.
- Eugene Vorontsov, Alican Bozkurt, Adam Casson, George Shaikovski, Michal Zelechowski, Kristen Severson, Eric Zimmermann, et al. A foundation model for clinical-grade computational pathology and rare cancers detection. *Nature Medicine*, 30(10):2924–2935, 2024. doi: 10.1038/s41591-024-03141-0.
- Eric Zimmermann, Eugene Vorontsov, Julian Viret, Adam Casson, Michal Zelechowski, and George Shaikovski. Virchow2. *2024 Shell TechXplorer Digest*, 2024(1), 2024. doi: 10.52196/208316.24.1.

A Full Calibration Metrics

Method	λ	NLL	NLL+TS	Brier	Brier+TS	ECE	ECE+TS	T
Balanced sampling	$\lambda = 0.5$	1.314 ± 0.117	0.839 ± 0.039	0.390 ± 0.013	0.339 ± 0.012	0.156 ± 0.005	0.016 ± 0.002	2.441 ± 0.095
Balanced sampling	$\lambda = 1.0$	1.301 ± 0.051	0.841 ± 0.051	0.391 ± 0.033	0.339 ± 0.026	0.158 ± 0.013	0.014 ± 0.003	2.334 ± 0.045
Balanced sampling	$\lambda = 1.5$	1.620 ± 0.112	1.060 ± 0.048	0.457 ± 0.023	0.399 ± 0.013	0.182 ± 0.011	0.017 ± 0.004	2.334 ± 0.045
CE + soft F1	$\lambda = 0.5$	1.626 ± 0.137	0.879 ± 0.044	0.413 ± 0.015	0.343 ± 0.011	0.179 ± 0.007	0.016 ± 0.004	2.884 ± 0.112
CE + soft F1	$\lambda = 1.0$	1.819 ± 0.050	0.889 ± 0.050	0.417 ± 0.032	0.343 ± 0.024	0.184 ± 0.012	0.014 ± 0.002	3.258 ± 0.124
CE + soft F1	$\lambda = 1.5$	2.274 ± 0.170	1.047 ± 0.060	0.486 ± 0.038	0.391 ± 0.029	0.217 ± 0.017	0.013 ± 0.003	3.520 ± 0.067
CE + soft MCC	$\lambda = 0.5$	1.375 ± 0.041	0.846 ± 0.017	0.397 ± 0.009	0.341 ± 0.009	0.162 ± 0.001	0.012 ± 0.004	2.552 ± 0.085
CE + soft MCC	$\lambda = 1.0$	1.392 ± 0.045	0.874 ± 0.039	0.411 ± 0.026	0.353 ± 0.023	0.169 ± 0.008	0.013 ± 0.002	2.496 ± 0.095
CE + soft MCC	$\lambda = 1.5$	1.713 ± 0.148	1.053 ± 0.059	0.451 ± 0.041	0.392 ± 0.030	0.183 ± 0.019	0.015 ± 0.005	2.497 ± 0.128
CFAL	$\lambda = 0.5$	3.232 ± 0.006	0.995 ± 0.042	0.952 ± 0.000	0.352 ± 0.014	0.723 ± 0.009	0.067 ± 0.007	0.050 ± 0.000
CFAL	$\lambda = 1.0$	3.228 ± 0.011	0.997 ± 0.087	0.951 ± 0.001	0.349 ± 0.029	0.728 ± 0.021	0.070 ± 0.008	0.050 ± 0.000
CFAL	$\lambda = 1.5$	3.246 ± 0.002	1.165 ± 0.049	0.953 ± 0.000	0.394 ± 0.018	0.695 ± 0.013	0.077 ± 0.006	0.050 ± 0.000
Focal loss	$\lambda = 0.5$	0.933 ± 0.040	0.819 ± 0.025	0.355 ± 0.010	0.336 ± 0.010	0.098 ± 0.003	0.010 ± 0.001	1.565 ± 0.030
Focal loss	$\lambda = 1.0$	0.933 ± 0.055	0.818 ± 0.053	0.357 ± 0.031	0.337 ± 0.027	0.100 ± 0.012	0.012 ± 0.002	1.548 ± 0.052
Focal loss	$\lambda = 1.5$	1.315 ± 0.076	0.970 ± 0.045	0.417 ± 0.017	0.373 ± 0.015	0.150 ± 0.008	0.013 ± 0.004	2.020 ± 0.000
OKO	$\lambda = 0.5$	0.931 ± 0.029	0.886 ± 0.031	0.352 ± 0.012	0.343 ± 0.012	0.071 ± 0.011	0.030 ± 0.003	0.804 ± 0.016
OKO	$\lambda = 1.0$	1.083 ± 0.076	0.902 ± 0.074	0.383 ± 0.028	0.337 ± 0.029	0.193 ± 0.006	0.041 ± 0.004	0.636 ± 0.012
OKO	$\lambda = 1.5$	1.188 ± 0.040	1.019 ± 0.047	0.410 ± 0.019	0.365 ± 0.022	0.192 ± 0.013	0.037 ± 0.007	0.651 ± 0.000
Weighted CE	$\lambda = 0.5$	1.303 ± 0.074	0.830 ± 0.027	0.387 ± 0.012	0.335 ± 0.011	0.156 ± 0.003	0.015 ± 0.002	2.441 ± 0.095
Weighted CE	$\lambda = 1.0$	1.297 ± 0.078	0.825 ± 0.056	0.387 ± 0.034	0.334 ± 0.028	0.157 ± 0.013	0.016 ± 0.005	2.413 ± 0.047
Weighted CE	$\lambda = 1.5$	1.530 ± 0.086	0.969 ± 0.043	0.431 ± 0.016	0.371 ± 0.015	0.174 ± 0.008	0.011 ± 0.006	2.440 ± 0.047

Table 7: Patch-feature calibration: NLL, Brier score, and ECE before (raw) and after (TS) temperature scaling, with fitted temperature T . One row per method per imbalance parameter.

Method	λ	NLL	NLL+TS	Brier	Brier+TS	ECE	ECE+TS	T
MDE-MIL	$\lambda = 0.5$	0.620 ± 0.021	0.588 ± 0.021	0.226 ± 0.010	0.221 ± 0.010	0.051 ± 0.007	0.037 ± 0.005	0.813 ± 0.032
MDE-MIL	$\lambda = 1.0$	0.582 ± 0.029	0.580 ± 0.029	0.220 ± 0.010	0.220 ± 0.011	0.042 ± 0.006	0.034 ± 0.007	1.049 ± 0.020
MDE-MIL	$\lambda = 1.5$	0.741 ± 0.053	0.740 ± 0.052	0.276 ± 0.013	0.276 ± 0.013	0.044 ± 0.004	0.039 ± 0.010	1.015 ± 0.039
Balanced sampling	$\lambda = 0.5$	0.873 ± 0.054	0.611 ± 0.022	0.271 ± 0.017	0.245 ± 0.012	0.105 ± 0.010	0.025 ± 0.007	1.999 ± 0.101
Balanced sampling	$\lambda = 1.0$	0.850 ± 0.030	0.611 ± 0.030	0.268 ± 0.016	0.246 ± 0.012	0.099 ± 0.012	0.020 ± 0.004	2.021 ± 0.067
Balanced sampling	$\lambda = 1.5$	1.140 ± 0.249	0.790 ± 0.137	0.330 ± 0.050	0.300 ± 0.044	0.124 ± 0.020	0.027 ± 0.003	2.143 ± 0.223
Focal loss	$\lambda = 0.5$	0.602 ± 0.036	0.563 ± 0.026	0.233 ± 0.014	0.228 ± 0.012	0.049 ± 0.013	0.019 ± 0.004	1.310 ± 0.000
Focal loss	$\lambda = 1.0$	0.624 ± 0.062	0.566 ± 0.055	0.233 ± 0.027	0.228 ± 0.025	0.054 ± 0.008	0.016 ± 0.005	1.416 ± 0.027
Focal loss	$\lambda = 1.5$	0.902 ± 0.090	0.744 ± 0.057	0.309 ± 0.034	0.290 ± 0.026	0.105 ± 0.020	0.031 ± 0.004	1.654 ± 0.000
Weighted CE	$\lambda = 0.5$	0.848 ± 0.055	0.600 ± 0.024	0.263 ± 0.018	0.238 ± 0.012	0.099 ± 0.012	0.022 ± 0.007	2.022 ± 0.115
Weighted CE	$\lambda = 1.0$	0.807 ± 0.068	0.582 ± 0.048	0.251 ± 0.025	0.231 ± 0.022	0.092 ± 0.010	0.022 ± 0.014	2.020 ± 0.000
Weighted CE	$\lambda = 1.5$	0.940 ± 0.070	0.708 ± 0.046	0.302 ± 0.024	0.275 ± 0.020	0.109 ± 0.009	0.023 ± 0.003	1.912 ± 0.073
RankMix	$\lambda = 0.5$	0.624 ± 0.045	0.623 ± 0.044	0.240 ± 0.019	0.239 ± 0.018	0.025 ± 0.015	0.020 ± 0.009	1.015 ± 0.052
RankMix	$\lambda = 1.0$	0.573 ± 0.055	0.569 ± 0.051	0.230 ± 0.026	0.229 ± 0.026	0.026 ± 0.005	0.018 ± 0.005	1.097 ± 0.021
RankMix	$\lambda = 1.5$	0.706 ± 0.063	0.654 ± 0.045	0.264 ± 0.016	0.257 ± 0.015	0.065 ± 0.008	0.026 ± 0.004	1.370 ± 0.026
SC-MIL	$\lambda = 0.5$	0.773 ± 0.044	0.577 ± 0.017	0.262 ± 0.010	0.239 ± 0.004	0.096 ± 0.006	0.026 ± 0.005	1.913 ± 0.098
SC-MIL	$\lambda = 1.0$	0.765 ± 0.087	0.586 ± 0.059	0.259 ± 0.031	0.239 ± 0.025	0.093 ± 0.011	0.018 ± 0.004	1.848 ± 0.036
SC-MIL	$\lambda = 1.5$	0.934 ± 0.070	0.693 ± 0.055	0.300 ± 0.012	0.274 ± 0.015	0.110 ± 0.001	0.024 ± 0.006	1.932 ± 0.037

Table 8: WSI-bag calibration: NLL, Brier score, and ECE before and after temperature scaling. Columns as in Table 7.